

Exploration_notebook

November 10, 2025

1 Predicting property sale prices in France: a study based on the DVF dataset

1.1 Exploratory notebook

1.2 Necessary libraires import

```
[2]: # Importing core libraries for data manipulation and visualization
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
import missingno as msno
import cartopy.crs as ccrs
import cartopy.feature as cfeature
import geopandas as gpd
import imageio.v2 as iio
import matplotlib.dates as mdates

# Plotly (Advanced graphs library) configuration
import plotly.io as pio
pio.renderers.default = "notebook"

# Warning management (to keep output clean)
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=UserWarning)
```

1.2.1 Load DVF Dataset

```
[3]: # Load the DVF dataset as DataFrame
df_dvf = pd.read_csv('dvf.csv', low_memory=False)

# Display the first 10 rows
df_dvf.head(10)
```

```
[3]: id_mutation date_mutation numero_disposition nature_mutation \
0    2020-1    2020-01-07           1      Vente
1    2020-2    2020-01-02           1      Vente
2    2020-2    2020-01-02           1      Vente
3    2020-2    2020-01-02           1      Vente
4    2020-2    2020-01-02           1      Vente
5    2020-2    2020-01-02           1      Vente
6    2020-3    2020-01-07           1      Vente
7    2020-4    2020-01-07           1      Vente
8    2020-4    2020-01-07           1      Vente
9    2020-5    2020-01-09           1      Vente

valeur_fonciere adresse_numero adresse_suffixe      adresse_nom_voie \
0        8000.0          NaN          NaN      FORTUNAT
1       2175.0          NaN          NaN  TERRES DES CINQ SAULES
2       2175.0          NaN          NaN     BOIS DU CHAMP RION
3       2175.0          NaN          NaN   EN COROBERT
4       2175.0          NaN          NaN  TERRES DES CINQ SAULES
5       2175.0          NaN          NaN  TERRES DES CINQ SAULES
6      75000.0          NaN          NaN  RUE DE LA CHARTREUSE
7        123.0          NaN          NaN    CHAMP PORTIER
8        123.0          NaN          NaN    CHAMP PORTIER
9      72000.0          NaN          NaN      CHAMPEL

adresse_code_voie code_postal ... type_local surface_reelle_bati \
0         B063     1250.0 ...      NaN          NaN
1         B124     1290.0 ...      NaN          NaN
2         B006     1290.0 ...      NaN          NaN
3         B025     1290.0 ...      NaN          NaN
4         B124     1290.0 ...      NaN          NaN
5         B124     1290.0 ...      NaN          NaN
6         0064     1960.0 ...      NaN          NaN
7         B041     1370.0 ...      NaN          NaN
8         B041     1370.0 ...      NaN          NaN
9         B034     1270.0 ...      NaN          NaN

nombre_pieces_principales code_nature_culture      nature_culture \
0             NaN              T      terres
1             NaN              BT  taillis simples
2             NaN              T      terres
3             NaN              T      terres
4             NaN              T      terres
5             NaN              T      terres
6             NaN              AB terrains a bâtir
7             NaN              S      sols
8             NaN              S      sols
9             NaN              J      jardins
```

```

code_nature_culture_speciale nature_culture_speciale surface_terrain \
0                      NaN                  NaN      1061.0
1                      NaN                  NaN       85.0
2                      NaN                  NaN     1115.0
3                      NaN                  NaN     1940.0
4                      NaN                  NaN     1148.0
5                      NaN                  NaN    2960.0
6                      NaN                  NaN      610.0
7                      NaN                  NaN       55.0
8                      NaN                  NaN      68.0
9                      NaN                  NaN      328.0

   longitude    latitude
0  5.323532  46.171941
1  4.893454  46.251858
2  4.900210  46.235277
3  4.882112  46.246554
4  4.894481  46.251841
5  4.894616  46.251941
6  5.226216  46.184570
7  5.344427  46.263955
8  5.343896  46.263803
9  5.350547  46.380898

[10 rows x 40 columns]

```

1.2.2 Inspect Dataset properties

Thanks to the **info** and **describe** methods, it is possible to quickly view the data contained in our dataset (entire CSV equivalent to 3GB). - The first method allows us to access the column names and their associated value types. - The second method allows us to view the orders of magnitude of the different columns.

```
[3]: # Check data types, non-null counts, and overall structure
df_dvf.info()

# Summary statistics for numerical columns (mean, std, min, max, and quartiles
# to understand data distribution)
df_dvf.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20133668 entries, 0 to 20133667
Data columns (total 40 columns):
 #   Column           Dtype  
 ---  --  
 0   id_mutation     object 
 1   date_mutation   object 
```

```

2    numero_disposition           int64
3    nature_mutation              object
4    valeur_fonciere               float64
5    adresse_numero                float64
6    adresse_suffixe               object
7    adresse_nom_voie              object
8    adresse_code_voie              object
9    code_postal                   float64
10   code_commune                  object
11   nom_commune                  object
12   code_departement              object
13   ancien_code_commune          float64
14   ancien_nom_commune           object
15   id_parcelle                  object
16   ancien_id_parcelle           object
17   numero_volume                 object
18   lot1_numero                  object
19   lot1_surface_carrez          float64
20   lot2_numero                  object
21   lot2_surface_carrez          float64
22   lot3_numero                  object
23   lot3_surface_carrez          float64
24   lot4_numero                  object
25   lot4_surface_carrez          float64
26   lot5_numero                  object
27   lot5_surface_carrez          float64
28   nombre_lots                  int64
29   code_type_local               float64
30   type_local                    object
31   surface_reelle_bati          float64
32   nombre_pieces_principales    float64
33   code_nature_culture           object
34   nature_culture                object
35   code_nature_culture_speciale object
36   nature_culture_speciale      object
37   surface_terrain                float64
38   longitude                     float64
39   latitude                      float64
dtypes: float64(15), int64(2), object(23)
memory usage: 6.0+ GB

```

```
[3]:      numero_disposition  valeur_fonciere  adresse_numero  code_postal  \
count      2.013367e+07    1.993772e+07    1.260362e+07  1.998002e+07
mean       1.234985e+00    1.529514e+06    7.150326e+02  5.000524e+04
std        7.899619e+00    1.684699e+07    2.017854e+03  2.739363e+04
min        1.000000e+00    1.000000e-02    1.000000e+00  1.000000e+03
25%       1.000000e+00    6.900000e+04    8.000000e+00  2.823000e+04
```

50%	1.000000e+00	1.650000e+05	2.500000e+01	4.914000e+04
75%	1.000000e+00	3.030000e+05	1.000000e+02	7.500700e+04
max	1.246000e+03	1.415000e+10	9.999000e+03	9.749000e+04

	ancien_code_commune	lot1_surface_carrez	lot2_surface_carrez	\
count	408.000000	1.761506e+06	557955.000000	
mean	20094.698529	6.735296e+01	64.035822	
std	18521.130057	2.082315e+02	71.815365	
min	14666.000000	1.000000e-02	0.010000	
25%	14666.000000	3.572000e+01	43.820000	
50%	14666.000000	5.510000e+01	61.570000	
75%	15031.000000	7.410000e+01	76.620000	
max	85212.000000	9.614000e+03	8705.000000	

	lot3_surface_carrez	lot4_surface_carrez	lot5_surface_carrez	\
count	61941.000000	14984.000000	5271.000000	
mean	71.970848	84.467058	96.111413	
std	98.255307	142.936682	196.619148	
min	0.200000	0.340000	0.400000	
25%	41.680000	39.617500	35.635000	
50%	62.080000	67.360000	70.080000	
75%	85.050000	100.370000	114.100000	
max	6947.850000	6947.850000	6947.850000	

	nombre_lots	code_type_local	surface_reelle_bati	\
count	2.013367e+07	1.183813e+07	6.983283e+06	
mean	4.320833e-01	2.226831e+00	1.153680e+02	
std	8.386084e-01	9.308073e-01	8.281463e+02	
min	0.000000e+00	1.000000e+00	1.000000e+00	
25%	0.000000e+00	1.000000e+00	5.000000e+01	
50%	0.000000e+00	2.000000e+00	7.500000e+01	
75%	1.000000e+00	3.000000e+00	1.050000e+02	
max	2.360000e+02	4.000000e+00	5.934000e+05	

	nombre_pieces_principales	surface_terrain	longitude	latitude
count	1.182539e+07	1.372763e+07	1.972417e+07	1.972417e+07
mean	1.863278e+00	2.863297e+03	2.354309e+00	4.610026e+01
std	2.091897e+00	1.409837e+04	6.379150e+00	5.921900e+00
min	0.000000e+00	1.000000e+00	-6.315108e+01	-2.138654e+01
25%	0.000000e+00	2.460000e+02	3.268020e-01	4.466372e+01
50%	1.000000e+00	6.220000e+02	2.365491e+00	4.672001e+01
75%	4.000000e+00	1.782000e+03	4.643447e+00	4.866416e+01
max	1.980000e+02	1.072309e+07	5.583079e+01	5.108645e+01

[4]: `print(df_dvf.columns) # Print dataframe columns names`


```
Index(['id_mutation', 'date_mutation', 'numero_disposition', 'nature_mutation',
       'valeur_fonciere', 'adresse_numero', 'adresse_suffixe',
```

```
'adresse_nom_voie', 'adresse_code_voie', 'code_postal', 'code_commune',
'nom_commune', 'code_departement', 'ancien_code_commune',
'ancien_nom_commune', 'id_parcelle', 'ancien_id_parcelle',
'numero_volume', 'lot1_numero', 'lot1_surface_carrez', 'lot2_numero',
'lot2_surface_carrez', 'lot3_numero', 'lot3_surface_carrez',
'lot4_numero', 'lot4_surface_carrez', 'lot5_numero',
'lot5_surface_carrez', 'nombre_lots', 'code_type_local', 'type_local',
'surface_reelle_bati', 'nombre_pieces_principales',
'code_nature_culture', 'nature_culture', 'code_nature_culture_speciale',
'nature_culture_speciale', 'surface_terrain', 'longitude', 'latitude'],
dtype='object')
```

1.2.3 Missing values overview

Thanks to the **missingno** library, it is possible to quickly view the null values contained in the columns as well as their total number. This allows us to easily identify whether it will be possible to delete columns or whether it will be necessary to evaluate this missing data.

```
[76]: msno.matrix(df_dvf)
plt.show()
```



```
[6]: df_dvf.isnull().sum() / len(df_dvf) * 100 # Calculate the percentage of missing data for each column
```

id_mutation	0.000000
date_mutation	0.000000
numero_disposition	0.000000
nature_mutation	0.000000
valeur_fonciere	0.973255

```

adresse_numero          37.400254
adresse_suffixe          95.517046
adresse_nom_voie        0.761684
adresse_code_voie        0.757830
code_postal              0.763164
code_commune             0.000000
nom_commune              0.000000
code_departement         0.000000
ancien_code_commune     99.997974
ancien_nom_commune      99.997974
id_parcelle              0.000000
ancien_id_parcelle       99.999930
numero_volume            99.770166
lot1_numero              68.647536
lot1_surface_carrez     91.250943
lot2_numero              91.009393
lot2_surface_carrez     97.228746
lot3_numero              98.398518
lot3_surface_carrez     99.692351
lot4_numero              99.486571
lot4_surface_carrez     99.925577
lot5_numero              99.779335
lot5_surface_carrez     99.973820
nombre_lots              0.000000
code_type_local           41.202304
type_local                41.202304
surface_reelle_bati      65.315396
nombre_pieces_principales 41.265610
code_nature_culture       31.814173
nature_culture             31.814173
code_nature_culture_speciale 95.707702
nature_culture_speciale   95.707702
surface_terrain            31.817556
longitude                  2.033917
latitude                   2.033917
dtype: float64

```

```
[7]: missing_percent = (df_dvf.isnull().sum() / len(df_dvf) * 100).
    ↪sort_values(ascending=False)

# Take only the 15 columns with more missing values
top_missing = missing_percent.head(20)


# Barplot
plt.figure(figsize=(10,6))
top_missing.plot(kind='bar', color='steelblue', edgecolor='black')
```

```

plt.title('Top 20 missing values columns', fontsize=14)
plt.ylabel('Percentage (%)', fontsize=12)
plt.xlabel('Columns', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()

```



To directly access the names of the ‘deletable’ columns, i.e. those with more than 90% missing values, we save them in a list.

```
[8]: # Delete columns
df_dvf = df_dvf.loc[:, df_dvf.isnull().mean() < 0.9]

df_dvf.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20133668 entries, 0 to 20133667
Data columns (total 24 columns):
 #   Column           Dtype  
 --- 
 0   id_mutation      object  
 1   date_mutation    object  
 2   numero_disposition   int64  
 3   nature_mutation   object  

```

```

4    valeur_fonciere           float64
5    adresse_numero            float64
6    adresse_nom_voie          object
7    adresse_code_voie          object
8    code_postal                float64
9    code_commune               object
10   nom_commune               object
11   code_departement          object
12   id_parcelle               object
13   lot1_numero               object
14   nombre_lots                int64
15   code_type_local            float64
16   type_local                 object
17   surface_reelle_bati        float64
18   nombre_pieces_principales float64
19   code_nature_culture        object
20   nature_culture              object
21   surface_terrain             float64
22   longitude                  float64
23   latitude                   float64
dtypes: float64(9), int64(2), object(13)
memory usage: 3.6+ GB

```

2 Correlation matrix

The correlation matrix provides an overview of the linear relationships between the numerical variables in the DVF dataset. Correlation values range from -1 (negative relationship) to +1 (positive relationship).


Most variables show weak linear relationships, as most correlations are near zero. `valeur_fonciere` (property value) has no strong correlation with other features, suggesting that prices depend on multiple non-linear factors such as location and property type. The strongest correlation is negative (-0.84) between `code_type_local` and `nombre_pieces_principales`, meaning that property type strongly influences room count. Overall, the dataset shows low multicollinearity, good for modeling, but linear models may not fully capture price dynamics.

```
[9]: corr = df_dvf.corr(numeric_only=True)

# Visualize the correlation matrix
plt.figure(figsize=(10,8))
sns.heatmap(corr,
            cmap='coolwarm',
            annot=True,
            fmt=".2f",
            linewidths=0.5,
            cbar=True)

plt.title("Correlation matrix", fontsize=14)
```

```
plt.show()
```




2.1 Property Type Distribution

The dataset shows that most transactions involve outbuildings and houses, while apartments and commercial premises are less frequent. This reflects the predominance of residential property sales in the DVF dataset.

```
[22]: # Visualize how many transactions correspond to each property category/type
df_dvf['type_local'].value_counts().plot(kind='bar', figsize=(6,4),  
                                         color='steelblue')

plt.title("Number of transactions by property type")
plt.ylabel("Count")
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.show()
```




2.1.1 Distribution of Sales Prices

The distribution of property sale prices is highly skewed, with most transactions involving relatively low to medium values. A few extreme outliers correspond to high-value properties. The **logarithmic scale** highlights the strong asymmetry typical of real estate markets.

```
[19]: # Plot histogram
plt.hist(df_dvf['valeur_fonciere'], bins=100, log=True)

plt.xscale('log')
plt.xlabel("Value (€)")
plt.ylabel("Count (log scale)")
plt.title("Distribution of sales prices")
```

```
plt.show()
```




2.1.2 Price Distribution by Property Type

The boxplot shows that property prices vary widely across all categories, with numerous high-value outliers. Houses and apartments have comparable median prices, while commercial and industrial properties show higher variability. The strong skew confirms the presence of extreme values.

```
[20]: # Use boxplots to compare the sale price ranges across property categories
df_dvf.boxplot(column='valeur_fonciere', by='type_local')

plt.yscale('log') # Log scale to handle extreme price variations
plt.ylabel("Value (€) [log]")
plt.title("Price distribution by property type")
plt.suptitle("") # Remove default title
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.show()
```



2.1.3 Seasonal plot

The number of transactions follows a clear seasonal pattern, typically peaking in summer and at the end of the year. Activity dropped significantly in 2020 probably due to the COVID-19 pandemic but gradually recovered in the following years


```
[17]: # Extract year and month
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'], errors='coerce')
df_dvf['year'] = df_dvf['date_mutation'].dt.year
df_dvf['month'] = df_dvf['date_mutation'].dt.month
```

```

# Count the number of transactions by (year, month)
monthly_sales = df_dvf.groupby(['year', 'month']).size().
    reset_index(name='count')

# Line plots: seasonality separated by year
plt.figure(figsize=(12,6))
sns.lineplot(
    data=monthly_sales,
    x='month', y='count', hue='year', palette='Spectral'
)
plt.title("Seasonal evolution of property sales per year")
plt.xlabel("Month")
plt.ylabel("Number of transactions")
plt.xticks(range(1,13))
plt.grid(alpha=0.3)
plt.show()

```



2.1.4 Impact of Surface Area on Property Prices


The scatter plot shows a positive relationship between surface area and property value, though with high dispersion. Larger properties generally sell for higher prices, but type differences and location effects create substantial variability.

```
[28]: # Scatter plot to visualize how surface area relates to property price
sns.scatterplot(
    data=df_dvf.sample(50000),
```

```

x='surface_reelle_bati', y='valeur_fonciere', hue='type_local', alpha=0.4
)
plt.xscale('log'); plt.yscale('log')
plt.title("Area vs Land value") #log scale
plt.xlabel("Surface (m2"); plt.ylabel("Valore (€)")
plt.show()

```



2.1.5 Average Monthly Sale Price Over Time

The average property value shows irregular fluctuations over time, with several extreme spikes likely caused by outlier transactions. These anomalies suggest the need for data cleaning or outlier filtering before modeling price trends.

```
[43]: # To track market changes over time, convert date to time and total mean sale
      ↵prices by month.
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'], ↵
                                         errors='coerce')
df_dvf['year_month'] = df_dvf['date_mutation'].dt.to_period('M')
```


```

monthly_price = (
    df_dvf.groupby('year_month')['valeur_fonciere']
    .mean()
    .reset_index()
)
monthly_price['year_month'] = monthly_price['year_month'].astype(str)

# Line plot showing monthly average price evolution
plt.figure(figsize=(12,8))
sns.lineplot(data=monthly_price, x='year_month', y='valeur_fonciere')

plt.title("Average monthly price over time")
plt.ylabel("Average value (€)")
plt.xticks(rotation=45, ha='right', fontsize=8)
plt.show()

```



2.1.6 Geographical Distribution of Property Sales in France

The geographical visualization highlights clear spatial patterns in real estate transactions across France. Most transactions are concentrated in Île-de-France, Auvergne-Rhône-Alpes, and Provence-Alpes-Côte d'Azur, which correspond to the country's most urbanized regions. The Paris metropolitan area in particular shows both a high density of transactions and the presence of some of the

highest property values.

Coastal regions, especially along the French Riviera and the Atlantic coast, also display clusters of high-value properties, reflecting strong demand in touristic and luxury real estate markets. In contrast, central and rural regions exhibit lower transaction density and generally lower prices, consistent with smaller populations and less market activity.

Overall, the map confirms a strong geographical inequality in the French housing market, with major cities and coastal zones concentrating the most expensive transactions.

```
[6]: # Create a geographic subset and log-transform property values
df_geo = df_dvf.dropna(subset=['longitude', 'latitude', 'valeur_fonciere']).copy()
df_geo['valeur_fonciere_log'] = np.log10(df_geo['valeur_fonciere'])

# Get 10'000 samples from the dataset (to have a rapid overview)
df_geo = df_geo.sample(10000, random_state=42)

fig, axes = plt.subplots(1, 2, figsize=(18, 9),
                       subplot_kw={'projection': ccrs.PlateCarree()})
ax1, ax2 = axes

# LEFT subplot: scatter plot with map background
ax1.add_feature(cfeature.BORDERS, linestyle='-', linewidth=0.5)
ax1.add_feature(cfeature.COASTLINE, linewidth=0.7)
ax1.add_feature(cfeature.LAND, edgecolor='black', facecolor='lightgray')
ax1.add_feature(cfeature.OCEAN, facecolor='lightblue')
ax1.set_extent([-5, 9, 41, 51], crs=ccrs.PlateCarree())

# Scatter plot
sc = ax1.scatter(
    df_geo['longitude'],
    df_geo['latitude'],
    c=df_geo['valeur_fonciere_log'],
    cmap='inferno',
    s=15,
    alpha=0.6,
    edgecolor='none',
    transform=ccrs.PlateCarree())
)

# Add colorbar to the first subplot (to differentiate values)
cbar = fig.colorbar(sc, ax=ax1, fraction=0.046, pad=0.04)
cbar.set_label('property value', fontsize=11)

# Add plot title & labels
ax1.set_title('Geographical distribution of real estate sales in France (Sample ↴10K)', fontsize=13)
```

```

ax1.set_xlabel('Longitude')
ax1.set_ylabel('Latitude')


# RIGHT subplot: map image
fig.delaxes(ax2) # remove the second cartopy subplot

ax_img = fig.add_axes([0.55, 0.1, 0.42, 0.8])

# Read and display the image
img = iio.imread("../docs/img/France_départementale.png")
ax_img.imshow(img)
ax_img.axis('off')
ax_img.set_title("France - Regions", fontsize=13)

# Adjust layout
plt.subplots_adjust(left=0.05, right=0.95, wspace=0.15)
plt.show()

```



2.1.7 Top 20 Departments by Number of Property Sales

The departments with the highest number of real estate transactions include Nord (59), Gironde (33), and Bouches-du-Rhône (13) — all of which are densely populated and economically active areas. These results align with national demographic patterns, as larger metropolitan regions tend to generate more sales activity. Departments such as Rhône (69), Alpes-Maritimes (06), and Haute-Garonne (31) also appear among the top 10, confirming strong real estate demand around major urban centers like Lyon, Nice, and Toulouse.

Overall, the chart highlights how property market activity is heavily concentrated in France's most urbanized regions.

```
[38]: # Extract department list with relative code
dept_list = df_dvf[['code_departement', 'nom_commune']].drop_duplicates() .
    ↪sort_values('code_departement')


# Print list
print(len(dept_list))
dept_list.head(10)
```

33378

```
[38]:   code_departement      nom_commune
0            01        Ceyzériat
13130         01        Ruffieu
12926         01     Échallon
12892         01     Mérignat
12786         01       Armix
12564         01  Boyeux-Saint-Jérôme
12552         01  Géovreisset
12478         01       Villes
12423         01      Conand
12204         01      Plagne
```

```
[77]: # Number of properties sold by department
sales_by_dept = (
    df_dvf.groupby('code_departement')
    .size()
    .reset_index(name='num_vendite')
    .sort_values(by='num_vendite', ascending=False)
)

plt.figure(figsize=(12,6))
sns.barplot(
    data=sales_by_dept.head(20),
    x='code_departement', y='num_vendite',
    palette='viridis'
)
plt.title("Top 20 Departments by number of properties sold", fontsize=14)
plt.xlabel('Departement')
plt.ylabel('Number of sales')
plt.grid(axis='y', alpha=0.3)
plt.show()
```



```
[ ]: # Evolution of average prices over time

# the date must be in datetime format
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],  
                                         errors='coerce')

# Add column Year-Month
df_period['year_month'] = df_period['date_mutation'].dt.to_period('M')

# Calculate average monthly price
monthly_avg = (
    df_period.groupby('year_month')['valeur_fonciere']
    .mean()
    .reset_index()
    .sort_values(by='year_month')
)
monthly_avg['year_month'] = monthly_avg['year_month'].astype(str)

# Line plot to show price's evolution
plt.figure(figsize=(12,6))
sns.lineplot(
    data=monthly_avg,
    x='year_month', y='valeur_fonciere',
    marker='o', color='darkorange'
)
plt.title('Evolution of average monthly prices (2020-2024)', fontsize=14)
plt.xlabel('Month')
```

```

plt.ylabel('Average price')
plt.xticks(rotation=45)
plt.grid(alpha=0.4)
plt.tight_layout()
plt.show()

```

```

-----
NameError                                 Traceback (most recent call last)
Cell In[7], line 7
      4 df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'], u
      ↪errors='coerce')
      6 # Add column Year-Month
----> 7 df_period['year_month'] = df_period['date_mutation'].dt.to_period('M')
      9 # Calculate average monthly price
     10 monthly_avg = (
     11     df_period.groupby('year_month')['valeur_fonciere']
     12     .mean()
     13     .reset_index()
     14     .sort_values(by='year_month')
     15 )

```


NameError: name 'df_period' is not defined

```

[78]: # Average prices by city
price_by_city = df_dvf.groupby('nom_commune')['valeur_fonciere'].mean().
    ↪reset_index()
top_cities = price_by_city.sort_values('valeur_fonciere', ascending=False).
    ↪head(20)

plt.figure(figsize=(10,6))
sns.barplot(data=top_cities, y='nom_commune', x='valeur_fonciere', u
    ↪palette='viridis')
plt.title("Top 20 municipalities by average sale price")
plt.xlabel("Average price")
plt.ylabel("City")
plt.show()

```



2.1.8 Evolution of average monthly prices (2020–2024)

The chart shows how the average property sale price in France has evolved month by month concerning the dataset period 2020 and 2024. Despite several sharp spikes likely caused by exceptional transactions or outliers, the overall trend appears relatively stable. This suggests that, while individual high-value sales occasionally distort the monthly average, the broader market did not experience major long-term volatility during this period.

```
[ ]: # Convert the date column to datetime format
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],  
                                         errors='coerce')

# Filter dataset for the years 2020-2024
df_period = df_dvf[df_dvf['date_mutation'].dt.year.between(2020, 2024)]

# Add Year-Month column for monthly aggregation
df_period['year_month'] = df_period['date_mutation'].dt.to_period('M')

# Calculate average sale price per month
monthly_avg = (
    df_period.groupby('year_month')['valeur_fonciere']
    .mean()
    .reset_index()
    .sort_values(by='year_month')
)

# Convert the period to datetime for proper time formatting
```

```

monthly_avg['year_month'] = monthly_avg['year_month'].dt.to_timestamp()


# Line plot showing temporal evolution of average prices
plt.figure(figsize=(12,6))
sns.lineplot(
    data=monthly_avg,
    x='year_month', y='valeur_fonciere',
    marker='o', color='darkorange'
)

plt.title('Evolution of average monthly prices', fontsize=14)
plt.xlabel('Month')
plt.ylabel('Average price (€)')

# X-axis: one major tick per year and minor ticks every 3 months
ax = plt.gca()
ax.xaxis.set_major_locator(mdates.YearLocator())
ax.xaxis.set_minor_locator(mdates.MonthLocator(interval=3))
ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))

plt.xticks(rotation=45)
plt.grid(alpha=0.4)
plt.tight_layout()
plt.show()

```




2.1.9 Number of transactions per month

The chart highlights the seasonality of real estate transactions in France. Sales activity tends to increase during the summer months, peaking in June and July, and again in December, likely due

to year-end transactions and fiscal timing. Conversely, August shows a notable drop, reflecting the slowdown typical of the vacation period.

```
[9]: # Extract the month from the transaction date and count how many sales occur in each month
df_dvf['month'] = df_dvf['date_mutation'].dt.month
transactions_per_month = df_dvf.groupby('month').size().
    reset_index(name='count')

# Bar plot showing the number of transactions by month
sns.barplot(
    data=transactions_per_month,
    x='month',
    y='count',
    palette='Blues_d'
)
plt.title("Number of transactions per month")
plt.xlabel("Month")
plt.ylabel("Number of sales")
plt.show()
```



2.2 Conclusion

This preliminary exploration of the DVF dataset provides a general understanding of its structure, data quality, and main patterns before conducting a formal Exploratory Data Analysis (EDA).

2.2.1 1. Dataset Structure and Quality

- The dataset is extremely large (~20 million records, 40 columns) and contains both administrative and transactional variables.
- Several cadastral-related columns show over 90% missing values and can be safely removed.
- The core variables relevant for predictive modeling — `valeur_fonciere`, `surface_reelle_bati`, `nombre_pieces_principales`, `type_local`, `date_mutation`, `longitude`, and `latitude` — are complete and consistent.

2.2.2 2. Property Value Distribution

- Property sale prices are heavily right-skewed, with a few extremely high values.
- A logarithmic transformation (\log_{10}) normalizes the distribution and improves interpretability.
- This confirms that future models should work on a log-transformed price variable.

2.2.3 3. Property Type Differences

- “Appartements” represent the majority of transactions, while “Maisons” tend to have higher median prices and larger variability.
- Property type is therefore a key categorical predictor to include in the Machine Learning models.

2.2.4 4. Relationship Between Area and Value

- A clear positive but sublinear correlation exists between built area and property value.
- Some small but high-priced properties likely correspond to urban premium locations.
- The price-per-square-meter ratio will be an important engineered feature.

2.2.5 5. Temporal Patterns

- Transaction volumes exhibit strong annual seasonality, with peaks in spring and early summer.
- Average monthly prices increase from 2020 to 2022, then stabilize slightly after 2023.
- Temporal variables should capture both trend and seasonality components.

2.2.6 6. Geographic Distribution

- Sales are concentrated in Île-de-France, the southeast coast, and major urban centers like the capital (Paris).
- Significant spatial heterogeneity is visible in both transaction density and price levels.
- Geographic variables (latitude, longitude, department) play a central role in property cost.

2.2.7 Summary

The DVF dataset provides a rich and reliable foundation for predictive modeling of property sale prices.

Despite some missing administrative fields, the essential features are clean and informative.

The observed temporal, spatial, and structural heterogeneity indicates that future models should combine log-transformed target values, spatial features, and time-based variables to capture the complex dynamics of the French real estate market.