

exploration

October 21, 2025

1 Predicting property sale prices in France: a study based on the DVF dataset

1.1 Exploratory notebook

1.2 Necessary libraires import

```
[2]: # Importing core libraries for data manipulation and visualization
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
import missingno as msno
import cartopy.crs as ccrs
import cartopy.feature as cfeature
import geopandas as gpd
import imageio.v2 as iio
import matplotlib.dates as mdates

# Plotly (Advanced graphs library) configuration
import plotly.io as pio
pio.renderers.default = "notebook"

# Warning management (to keep output clean)
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=UserWarning)
```

1.2.1 Load DVF Dataset

```
[3]: # Load the DVF dataset as DataFrame
df_dvf = pd.read_csv('dvf.csv', low_memory=False)

# Display the first 10 rows
df_dvf.head(10)
```

```

[3]:  id_mutation date_mutation  numero_disposition nature_mutation \
0      2020-1      2020-01-07                1      Vente
1      2020-2      2020-01-02                1      Vente
2      2020-2      2020-01-02                1      Vente
3      2020-2      2020-01-02                1      Vente
4      2020-2      2020-01-02                1      Vente
5      2020-2      2020-01-02                1      Vente
6      2020-3      2020-01-07                1      Vente
7      2020-4      2020-01-07                1      Vente
8      2020-4      2020-01-07                1      Vente
9      2020-5      2020-01-09                1      Vente

      valeur_fonciere  adresse_numero  adresse_suffixe      adresse_nom_voie \
0              8000.0              NaN              NaN      FORTUNAT
1              2175.0              NaN              NaN  TERRES DES CINQ SAULES
2              2175.0              NaN              NaN    BOIS DU CHAMP RION
3              2175.0              NaN              NaN      EN COROBERT
4              2175.0              NaN              NaN  TERRES DES CINQ SAULES
5              2175.0              NaN              NaN  TERRES DES CINQ SAULES
6             75000.0              NaN              NaN    RUE DE LA CHARTREUSE
7               123.0              NaN              NaN    CHAMP PORTIER
8               123.0              NaN              NaN    CHAMP PORTIER
9             72000.0              NaN              NaN    CHAMPEL

      adresse_code_voie  code_postal  ... type_local  surface_reelle_bati \
0              B063      1250.0  ...      NaN              NaN
1              B124      1290.0  ...      NaN              NaN
2              B006      1290.0  ...      NaN              NaN
3              B025      1290.0  ...      NaN              NaN
4              B124      1290.0  ...      NaN              NaN
5              B124      1290.0  ...      NaN              NaN
6              0064      1960.0  ...      NaN              NaN
7              B041      1370.0  ...      NaN              NaN
8              B041      1370.0  ...      NaN              NaN
9              B034      1270.0  ...      NaN              NaN

      nombre_pieces_principales  code_nature_culture      nature_culture \
0              NaN              T      terres
1              NaN              BT  taillis simples
2              NaN              T      terres
3              NaN              T      terres
4              NaN              T      terres
5              NaN              T      terres
6              NaN              AB  terrains a bâtir
7              NaN              S      sols
8              NaN              S      sols
9              NaN              J      jardins

```

| | code_nature_culture_speciale | nature_culture_speciale | surface_terrain | \ |
|---|------------------------------|-------------------------|-----------------|---|
| 0 | NaN | NaN | 1061.0 | |
| 1 | NaN | NaN | 85.0 | |
| 2 | NaN | NaN | 1115.0 | |
| 3 | NaN | NaN | 1940.0 | |
| 4 | NaN | NaN | 1148.0 | |
| 5 | NaN | NaN | 2960.0 | |
| 6 | NaN | NaN | 610.0 | |
| 7 | NaN | NaN | 55.0 | |
| 8 | NaN | NaN | 68.0 | |
| 9 | NaN | NaN | 328.0 | |

| | longitude | latitude |
|---|-----------|-----------|
| 0 | 5.323532 | 46.171941 |
| 1 | 4.893454 | 46.251858 |
| 2 | 4.900210 | 46.235277 |
| 3 | 4.882112 | 46.246554 |
| 4 | 4.894481 | 46.251841 |
| 5 | 4.894616 | 46.251941 |
| 6 | 5.226216 | 46.184570 |
| 7 | 5.344427 | 46.263955 |
| 8 | 5.343896 | 46.263803 |
| 9 | 5.350547 | 46.380898 |

[10 rows x 40 columns]

1.2.2 Inspect Dataset properties

Thanks to the **info** and **describe** methods, it is possible to quickly view the data contained in our dataset (entire CSV equivalent to 3GB). - The first method allows us to access the column names and their associated value types. - The second method allows us to view the orders of magnitude of the different columns.

```
[3]: # Check data types, non-null counts, and overall structure
df_dvf.info()

# Summary statistics for numerical columns (mean, std, min, max, and quartiles)
↳ to understand data distribution
df_dvf.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20133668 entries, 0 to 20133667
Data columns (total 40 columns):
#   Column                Dtype
---  -
0   id_mutation           object
1   date_mutation         object
```

```

2  numero_disposition      int64
3  nature_mutation         object
4  valeur_fonciere         float64
5  adresse_numero          float64
6  adresse_suffixe         object
7  adresse_nom_voie        object
8  adresse_code_voie       object
9  code_postal             float64
10 code_commune            object
11 nom_commune             object
12 code_departement        object
13 ancien_code_commune     float64
14 ancien_nom_commune      object
15 id_parcelle             object
16 ancien_id_parcelle      object
17 numero_volume           object
18 lot1_numero             object
19 lot1_surface_carrez     float64
20 lot2_numero             object
21 lot2_surface_carrez     float64
22 lot3_numero             object
23 lot3_surface_carrez     float64
24 lot4_numero             object
25 lot4_surface_carrez     float64
26 lot5_numero             object
27 lot5_surface_carrez     float64
28 nombre_lots             int64
29 code_type_local         float64
30 type_local              object
31 surface_reelle_bati     float64
32 nombre_pieces_principales float64
33 code_nature_culture     object
34 nature_culture          object
35 code_nature_culture_speciale object
36 nature_culture_speciale object
37 surface_terrain        float64
38 longitude               float64
39 latitude                float64
dtypes: float64(15), int64(2), object(23)
memory usage: 6.0+ GB

```

```

[3]:      numero_disposition  valeur_fonciere  adresse_numero  code_postal  \
count      2.013367e+07      1.993772e+07      1.260362e+07      1.998002e+07
mean       1.234985e+00      1.529514e+06      7.150326e+02      5.000524e+04
std        7.899619e+00      1.684699e+07      2.017854e+03      2.739363e+04
min        1.000000e+00      1.000000e-02      1.000000e+00      1.000000e+03
25%        1.000000e+00      6.900000e+04      8.000000e+00      2.823000e+04

```

| | | | | |
|-----|--------------|--------------|--------------|--------------|
| 50% | 1.000000e+00 | 1.650000e+05 | 2.500000e+01 | 4.914000e+04 |
| 75% | 1.000000e+00 | 3.030000e+05 | 1.000000e+02 | 7.500700e+04 |
| max | 1.246000e+03 | 1.415000e+10 | 9.999000e+03 | 9.749000e+04 |

| | ancien_code_commune | lot1_surface_carrez | lot2_surface_carrez | \ |
|-------|---------------------|---------------------|---------------------|---|
| count | 408.000000 | 1.761506e+06 | 557955.000000 | |
| mean | 20094.698529 | 6.735296e+01 | 64.035822 | |
| std | 18521.130057 | 2.082315e+02 | 71.815365 | |
| min | 14666.000000 | 1.000000e-02 | 0.010000 | |
| 25% | 14666.000000 | 3.572000e+01 | 43.820000 | |
| 50% | 14666.000000 | 5.510000e+01 | 61.570000 | |
| 75% | 15031.000000 | 7.410000e+01 | 76.620000 | |
| max | 85212.000000 | 9.614000e+03 | 8705.000000 | |

| | lot3_surface_carrez | lot4_surface_carrez | lot5_surface_carrez | \ |
|-------|---------------------|---------------------|---------------------|---|
| count | 61941.000000 | 14984.000000 | 5271.000000 | |
| mean | 71.970848 | 84.467058 | 96.111413 | |
| std | 98.255307 | 142.936682 | 196.619148 | |
| min | 0.200000 | 0.340000 | 0.400000 | |
| 25% | 41.680000 | 39.617500 | 35.635000 | |
| 50% | 62.080000 | 67.360000 | 70.080000 | |
| 75% | 85.050000 | 100.370000 | 114.100000 | |
| max | 6947.850000 | 6947.850000 | 6947.850000 | |

| | nombre_lots | code_type_local | surface_reelle_bati | \ |
|-------|--------------|-----------------|---------------------|---|
| count | 2.013367e+07 | 1.183813e+07 | 6.983283e+06 | |
| mean | 4.320833e-01 | 2.226831e+00 | 1.153680e+02 | |
| std | 8.386084e-01 | 9.308073e-01 | 8.281463e+02 | |
| min | 0.000000e+00 | 1.000000e+00 | 1.000000e+00 | |
| 25% | 0.000000e+00 | 1.000000e+00 | 5.000000e+01 | |
| 50% | 0.000000e+00 | 2.000000e+00 | 7.500000e+01 | |
| 75% | 1.000000e+00 | 3.000000e+00 | 1.050000e+02 | |
| max | 2.360000e+02 | 4.000000e+00 | 5.934000e+05 | |

| | nombre_pieces_principales | surface_terrain | longitude | latitude |
|-------|---------------------------|-----------------|---------------|---------------|
| count | 1.182539e+07 | 1.372763e+07 | 1.972417e+07 | 1.972417e+07 |
| mean | 1.863278e+00 | 2.863297e+03 | 2.354309e+00 | 4.610026e+01 |
| std | 2.091897e+00 | 1.409837e+04 | 6.379150e+00 | 5.921900e+00 |
| min | 0.000000e+00 | 1.000000e+00 | -6.315108e+01 | -2.138654e+01 |
| 25% | 0.000000e+00 | 2.460000e+02 | 3.268020e-01 | 4.466372e+01 |
| 50% | 1.000000e+00 | 6.220000e+02 | 2.365491e+00 | 4.672001e+01 |
| 75% | 4.000000e+00 | 1.782000e+03 | 4.643447e+00 | 4.866416e+01 |
| max | 1.980000e+02 | 1.072309e+07 | 5.583079e+01 | 5.108645e+01 |

```
[4]: print(df_dvf.columns) # Print dataframe columns names
```


```
Index(['id_mutation', 'date_mutation', 'numero_disposition', 'nature_mutation',  
      'valeur_fonciere', 'adresse_numero', 'adresse_suffixe',
```

```
'adresse_nom_voie', 'adresse_code_voie', 'code_postal', 'code_commune',
'nom_commune', 'code_departement', 'ancien_code_commune',
'ancien_nom_commune', 'id_parcelle', 'ancien_id_parcelle',
'numero_volume', 'lot1_numero', 'lot1_surface_carrez', 'lot2_numero',
'lot2_surface_carrez', 'lot3_numero', 'lot3_surface_carrez',
'lot4_numero', 'lot4_surface_carrez', 'lot5_numero',
'lot5_surface_carrez', 'nombre_lots', 'code_type_local', 'type_local',
'surface_reelle_bati', 'nombre_pieces_principales',
'code_nature_culture', 'nature_culture', 'code_nature_culture_speciale',
'nature_culture_speciale', 'surface_terrain', 'longitude', 'latitude'],
dtype='object')
```

1.2.3 Missing values overview

Thanks to the **missingno** library, it is possible to quickly view the null values contained in the columns as well as their total number. This allows us to easily identify whether it will be possible to delete columns or whether it will be necessary to evaluate this missing data.

```
[76]: msno.matrix(df_dvf)
plt.show()
```



```
[6]: df_dvf.isnull().sum() / len(df_dvf) * 100 # Calculate the percentage of missing
      ↪ data for each column
```

```
[6]: id_mutation          0.000000
     date_mutation        0.000000
     numero_disposition    0.000000
     nature_mutation       0.000000
     valeur_fonciere       0.973255
```

| | |
|------------------------------|-----------|
| adresse_numero | 37.400254 |
| adresse_suffixe | 95.517046 |
| adresse_nom_voie | 0.761684 |
| adresse_code_voie | 0.757830 |
| code_postal | 0.763164 |
| code_commune | 0.000000 |
| nom_commune | 0.000000 |
| code_departement | 0.000000 |
| ancien_code_commune | 99.997974 |
| ancien_nom_commune | 99.997974 |
| id_parcelle | 0.000000 |
| ancien_id_parcelle | 99.999930 |
| numero_volume | 99.770166 |
| lot1_numero | 68.647536 |
| lot1_surface_carrez | 91.250943 |
| lot2_numero | 91.009393 |
| lot2_surface_carrez | 97.228746 |
| lot3_numero | 98.398518 |
| lot3_surface_carrez | 99.692351 |
| lot4_numero | 99.486571 |
| lot4_surface_carrez | 99.925577 |
| lot5_numero | 99.779335 |
| lot5_surface_carrez | 99.973820 |
| nombre_lots | 0.000000 |
| code_type_local | 41.202304 |
| type_local | 41.202304 |
| surface_reelle_bati | 65.315396 |
| nombre_pieces_principales | 41.265610 |
| code_nature_culture | 31.814173 |
| nature_culture | 31.814173 |
| code_nature_culture_speciale | 95.707702 |
| nature_culture_speciale | 95.707702 |
| surface_terrain | 31.817556 |
| longitude | 2.033917 |
| latitude | 2.033917 |

dtype: float64

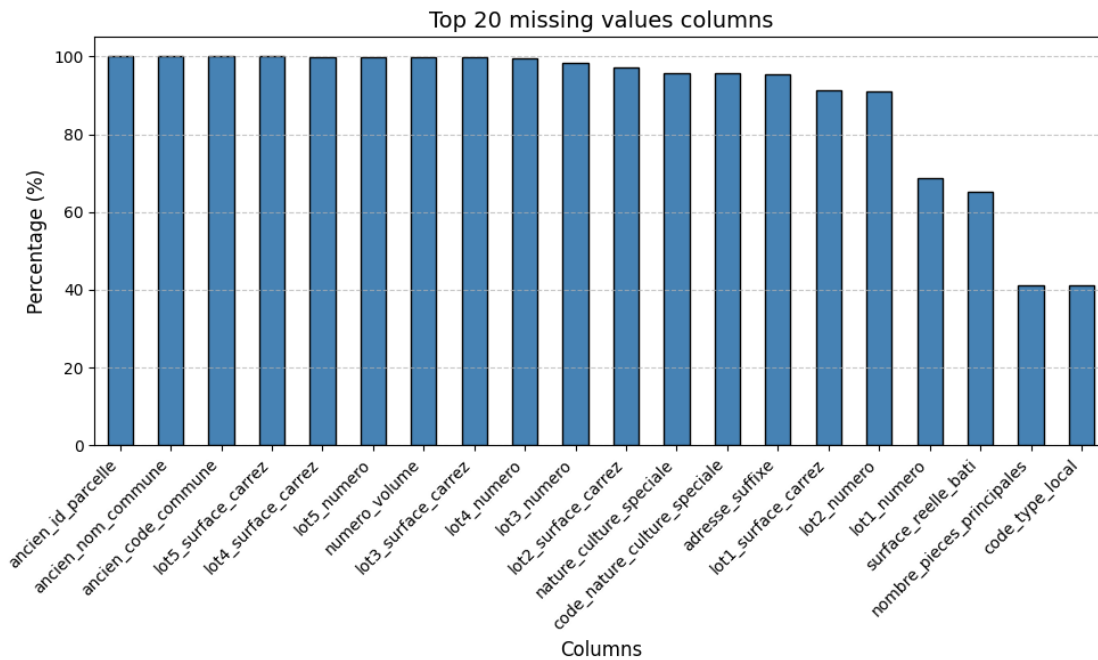
```
[7]: missing_percent = (df_dvf.isnull().sum() / len(df_dvf) * 100).
      ↪sort_values(ascending=False)

# Take only the 15 columns with more missing values
top_missing = missing_percent.head(20)

# Barplot
plt.figure(figsize=(10,6))
top_missing.plot(kind='bar', color='steelblue', edgecolor='black')
```

```
plt.title('Top 20 missing values columns', fontsize=14)
plt.ylabel('Percentage (%)', fontsize=12)
plt.xlabel('Columns', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()
```



To directly access the names of the ‘deletable’ columns, i.e. those with more than 90% missing values, we save them in a list.

```
[8]: # Delete columns
df_dvf = df_dvf.loc[:, df_dvf.isnull().mean() < 0.9]

df_dvf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20133668 entries, 0 to 20133667
Data columns (total 24 columns):
#   Column                Dtype
---  -
0   id_mutation           object
1   date_mutation         object
2   numero_disposition    int64
3   nature_mutation       object
```



```

4  valeur_fonciere          float64
5  adresse_numero          float64
6  adresse_nom_voie        object
7  adresse_code_voie        object
8  code_postal             float64
9  code_commune            object
10 nom_commune             object
11 code_departement        object
12 id_parcelle            object
13 lot1_numero            object
14 nombre_lots            int64
15 code_type_local         float64
16 type_local             object
17 surface_reelle_bati     float64
18 nombre_pieces_principales float64
19 code_nature_culture      object
20 nature_culture          object
21 surface_terrain        float64
22 longitude              float64
23 latitude               float64
dtypes: float64(9), int64(2), object(13)
memory usage: 3.6+ GB

```

2 Correlation matrix

The correlation matrix provides an overview of the linear relationships between the numerical variables in the DVF dataset. Correlation values range from -1 (negative relationship) to +1 (positive relationship).

Most variables show weak linear relationships, as most correlations are near zero. **valeur_fonciere** (property value) has no strong correlation with other features, suggesting that prices depend on multiple non-linear factors such as location and property type. The strongest correlation is negative (-0.84) between **code_type_local** and **nombre_pieces_principales**, meaning that property type strongly influences room count. Overall, the dataset shows low multicollinearity, good for modeling, but linear models may not fully capture price dynamics.

```

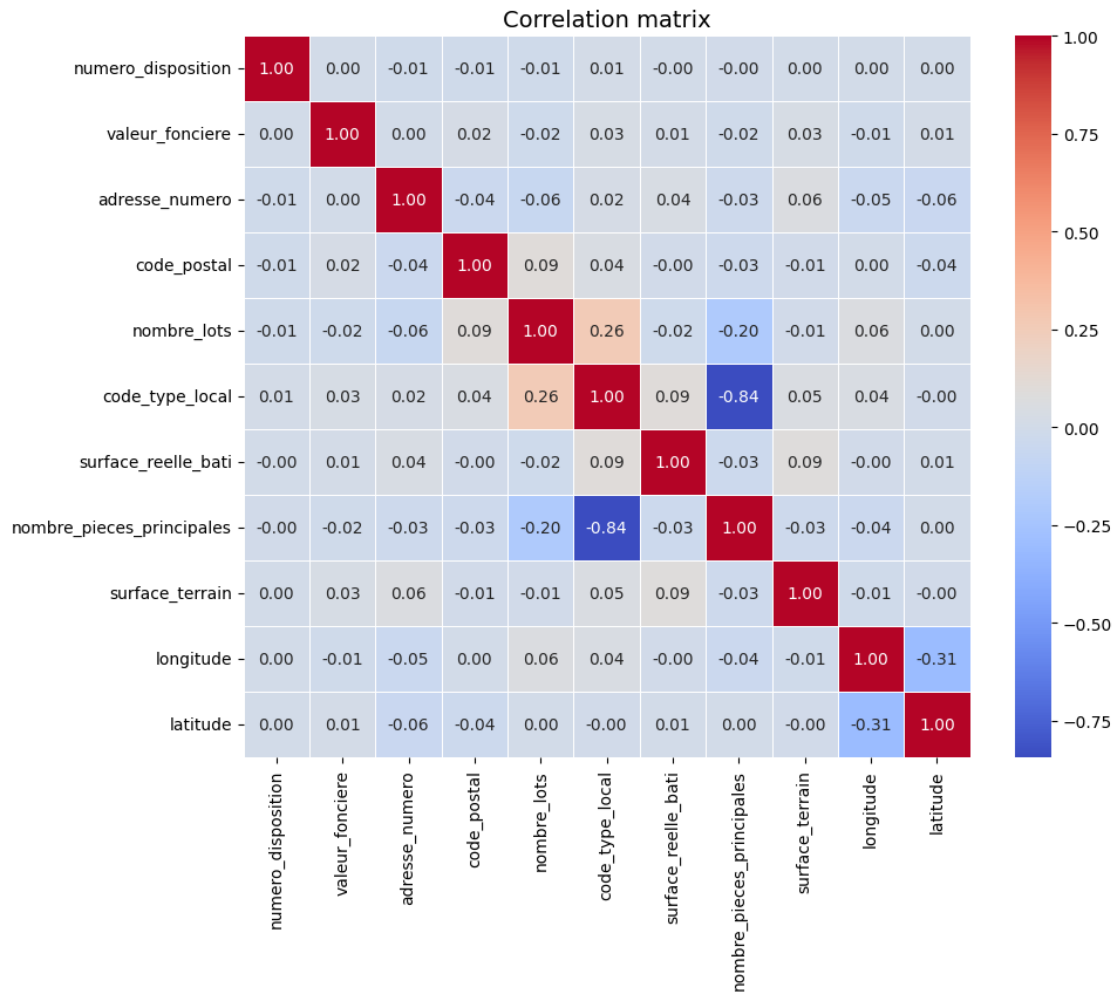
[9]: corr = df_dvf.corr(numeric_only=True)

# Visualize the correlation matrix
plt.figure(figsize=(10,8))
sns.heatmap(corr,
            cmap='coolwarm',
            annot=True,
            fmt=".2f",
            linewidths=0.5,
            cbar=True)

plt.title("Correlation matrix", fontsize=14)

```

```
plt.show()
```




2.1 Property Type Distribution

The dataset shows that most transactions involve outbuildings and houses, while apartments and commercial premises are less frequent. This reflects the predominance of residential property sales in the DVF dataset.

```
[22]: # Visualize how many transactions correspond to each property category/type
df_dvf['type_local'].value_counts().plot(kind='bar', figsize=(6,4),
color='steelblue')

plt.title("Number of transactions by property type")
plt.ylabel("Count")
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.show()
```




2.1.1 Distribution of Sales Prices

The distribution of property sale prices is highly skewed, with most transactions involving relatively low to medium values. A few extreme outliers correspond to high-value properties. The **logarithmic scale** highlights the strong asymmetry typical of real estate markets.

```
[19]: # Plot histogram
plt.hist(df_dvf['valeur_fonciere'], bins=100, log=True)

plt.xscale('log')
plt.xlabel("Value (€)")
plt.ylabel("Count (log scale)")
plt.title("Distribution of sales prices")
```

```
plt.show()
```




2.1.2 Price Distribution by Property Type

The boxplot shows that property prices vary widely across all categories, with numerous high-value outliers. Houses and apartments have comparable median prices, while commercial and industrial properties show higher variability. The strong skew confirms the presence of extreme values.

```
[20]: # Use boxplots to compare the sale price ranges across property categories
df_dvf.boxplot(column='valeur_fonciere', by='type_local')

plt.yscale('log') # Log scale to handle extreme price variations
plt.ylabel("Value (€) [log]")
plt.title("Price distribution by property type")
plt.suptitle("") # Remove default title
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.show()
```



2.1.3 Seasonal plot

The number of transactions follows a clear seasonal pattern, typically peaking in summer and at the end of the year. Activity dropped significantly in 2020 probably due to the COVID-19 pandemic but gradually recovered in the following years

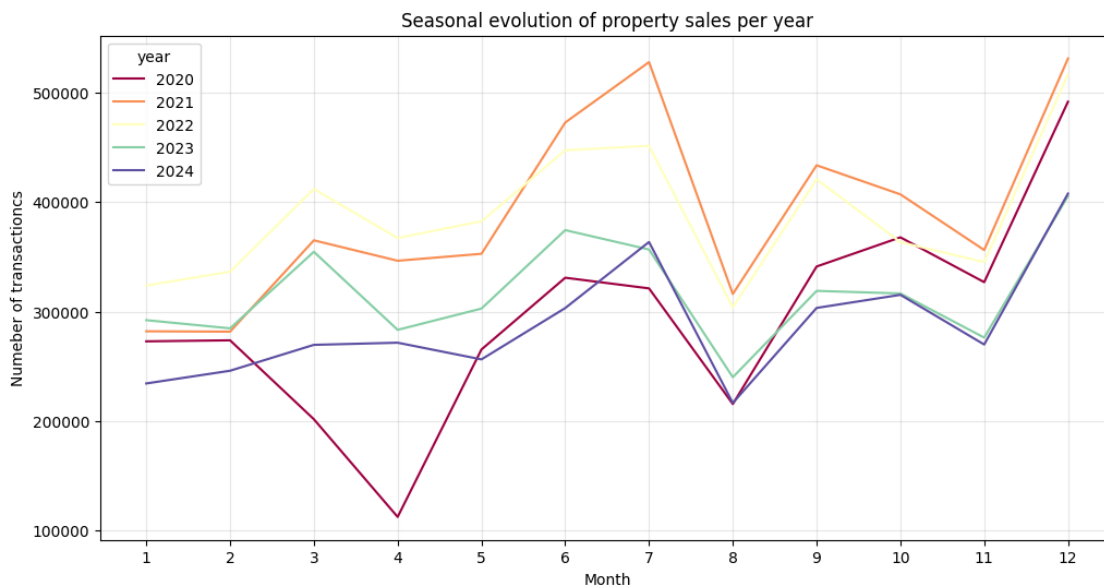
```
[17]: # Extract year and month
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],
    errors='coerce')
df_dvf['year'] = df_dvf['date_mutation'].dt.year
df_dvf['month'] = df_dvf['date_mutation'].dt.month
```

```

# Count the number of transactions by (year, month)
monthly_sales = df_dvf.groupby(['year', 'month']).size().
    ↪reset_index(name='count')

# Line plots: seasonality separated by year
plt.figure(figsize=(12,6))
sns.lineplot(
    data=monthly_sales,
    x='month', y='count', hue='year', palette='Spectral'
)
plt.title("Seasonal evolution of property sales per year")
plt.xlabel("Month")
plt.ylabel("Numeber of transactioncs")
plt.xticks(range(1,13))
plt.grid(alpha=0.3)
plt.show()

```



2.1.4 Impact of Surface Area on Property Prices

The scatter plot shows a positive relationship between surface area and property value, though with high dispersion. Larger properties generally sell for higher prices, but type differences and location effects create substantial variability.

```


[28]: # Scatter plot to visualize how surface area relates to property price
sns.scatterplot(
    data=df_dvf.sample(50000),

```

```

x='surface_reelle_bati', y='valeur_fonciere', hue='type_local', alpha=0.4
)
plt.xscale('log'); plt.yscale('log')
plt.title("Area vs Land value") #log scale
plt.xlabel("Surface (m²)"); plt.ylabel("Valore (€)")
plt.show()

```



2.1.5 Average Monthly Sale Price Over Time

The average property value shows irregular fluctuations over time, with several extreme spikes likely caused by outlier transactions. These anomalies suggest the need for data cleaning or outlier filtering before modeling price trends.

```

[43]: # To track market changes over time, convert date to time and total mean sale
      ↪ prices by month.
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],
      ↪ errors='coerce')
df_dvf['year_month'] = df_dvf['date_mutation'].dt.to_period('M')

```


```

monthly_price = (
    df_dvf.groupby('year_month')['valeur_fonciere']
    .mean()
    .reset_index()
)
monthly_price['year_month'] = monthly_price['year_month'].astype(str)

# Line plot showing monthly average price evolution
plt.figure(figsize=(12,8))
sns.lineplot(data=monthly_price, x='year_month', y='valeur_fonciere')

plt.title("Average monthly price over time")
plt.ylabel("Average value (€)")
plt.xticks(rotation=45, ha='right', fontsize=8)
plt.show()

```



2.1.6 Geographical Distribution of Property Sales in France

The geographical visualization highlights clear spatial patterns in real estate transactions across France. Most transactions are concentrated in Île-de-France, Auvergne-Rhône-Alpes, and Provence-Alpes-Côte d'Azur, which correspond to the country's most urbanized regions. The Paris metropolitan area in particular shows both a high density of transactions and the presence of some of the

highest property values.

Coastal regions, especially along the French Riviera and the Atlantic coast, also display clusters of high-value properties, reflecting strong demand in touristic and luxury real estate markets. In contrast, central and rural regions exhibit lower transaction density and generally lower prices, consistent with smaller populations and less market activity.

Overall, the map confirms a strong geographical inequality in the French housing market, with major cities and coastal zones concentrating the most expensive transactions.

```
[6]: # Create a geographic subset and log-transform property values
df_geo = df_dvf.dropna(subset=['longitude', 'latitude', 'valeur_fonciere']).
    ↪copy()
df_geo['valeur_fonciere_log'] = np.log10(df_geo['valeur_fonciere'])

# Get 10'000 samples from the dataset (to have a rapid overview)
df_geo = df_geo.sample(10000, random_state=42)

fig, axes = plt.subplots(1, 2, figsize=(18, 9),
                        subplot_kw={'projection': ccrs.PlateCarree()})
ax1, ax2 = axes

# LEFT subplot: scatter plot with map background
ax1.add_feature(cfeature.BORDERS, linestyle='-', linewidth=0.5)
ax1.add_feature(cfeature.COASTLINE, linewidth=0.7)
ax1.add_feature(cfeature.LAND, edgecolor='black', facecolor='lightgray')
ax1.add_feature(cfeature.OCEAN, facecolor='lightblue')
ax1.set_extent([-5, 9, 41, 51], crs=ccrs.PlateCarree())

# Scatter plot
sc = ax1.scatter(
    df_geo['longitude'],
    df_geo['latitude'],
    c=df_geo['valeur_fonciere_log'],
    cmap='inferno',
    s=15,
    alpha=0.6,
    edgecolor='none',
    transform=ccrs.PlateCarree()
)

# Add colorbar to the first subplot (to differentiate values)
cbar = fig.colorbar(sc, ax=ax1, fraction=0.046, pad=0.04)
cbar.set_label('property value', fontsize=11)

# Add plot title & labels
ax1.set_title('Geographical distribution of real estate sales in France (Sample_
    ↪10K)', fontsize=13)
```

```

ax1.set_xlabel('Longitude')
ax1.set_ylabel('Latitude')

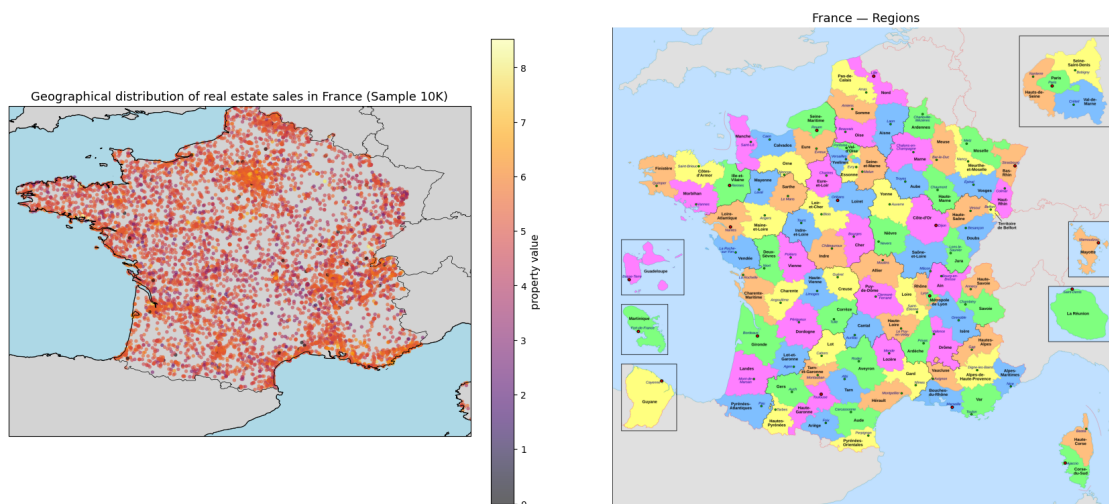
# RIGHT subplot: map image
fig.delaxes(ax2) # remove the second cartopy subplot

ax_img = fig.add_axes([0.55, 0.1, 0.42, 0.8])

# Read and display the image
img = io.imread("../docs/img/France_départementale.png")
ax_img.imshow(img)
ax_img.axis('off')
ax_img.set_title("France - Regions", fontsize=13)

# Adjust layout
plt.subplots_adjust(left=0.05, right=0.95, wspace=0.15)
plt.show()

```



2.1.7 Top 20 Departments by Number of Property Sales

The departments with the highest number of real estate transactions include Nord (59), Gironde (33), and Bouches-du-Rhône (13) — all of which are densely populated and economically active areas. These results align with national demographic patterns, as larger metropolitan regions tend to generate more sales activity. Departments such as Rhône (69), Alpes-Maritimes (06), and Haute-Garonne (31) also appear among the top 10, confirming strong real estate demand around major urban centers like Lyon, Nice, and Toulouse.

Overall, the chart highlights how property market activity is heavily concentrated in France's most urbanized regions.

```
[38]: # Extract deparment list with relative code
dept_list = df_dvf[['code_departement', 'nom_commune']].drop_duplicates().
↳sort_values('code_departement')

# Print list
print(len(dept_list))
dept_list.head(10)
```

33378


```
[38]:
```

| | code_departement | nom_commune |
|-------|------------------|---------------------|
| 0 | 01 | Ceyzériat |
| 13130 | 01 | Ruffieu |
| 12926 | 01 | Échallon |
| 12892 | 01 | Mérignat |
| 12786 | 01 | Armix |
| 12564 | 01 | Boyeux-Saint-Jérôme |
| 12552 | 01 | Géovreisset |
| 12478 | 01 | Villes |
| 12423 | 01 | Conand |
| 12204 | 01 | Plagne |

```
[77]: # Number of properties sold by department
sales_by_dept = (
    df_dvf.groupby('code_departement')
    .size()
    .reset_index(name='num_vendite')
    .sort_values(by='num_vendite', ascending=False)
)

plt.figure(figsize=(12,6))
sns.barplot(
    data=sales_by_dept.head(20),
    x='code_departement', y='num_vendite',
    palette='viridis'
)

plt.title("Top 20 Departments by number of properties sold", fontsize=14)
plt.xlabel('Departement')
plt.ylabel('Number of sales')
plt.grid(axis='y', alpha=0.3)
plt.show()
```



```
[7]: # Evolution of average prices over time

# the date must be in datetime format
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],
    ↪errors='coerce')

# Add column Year-Month
df_period['year_month'] = df_period['date_mutation'].dt.to_period('M')

# Calculate average monthly price
monthly_avg = (
    df_period.groupby('year_month')['valeur_fonciere']
    .mean()
    .reset_index()
    .sort_values(by='year_month')
)
monthly_avg['year_month'] = monthly_avg['year_month'].astype(str)

# Line plot to show price's evolution
plt.figure(figsize=(12,6))
sns.lineplot(
    data=monthly_avg,
    x='year_month', y='valeur_fonciere',
    marker='o', color='darkorange'
)
plt.title('Evolution of average monthly prices (2019-2024)', fontsize=14)
plt.xlabel('Month')
```


```
plt.ylabel('Average price')
plt.xticks(rotation=45)
plt.grid(alpha=0.4)
plt.tight_layout()
plt.show()
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[7], line 7
      4 df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],
      ↪errors='coerce')
      6 # Add column Year-Month
----> 7 df_period['year_month'] = df_period['date_mutation'].dt.to_period('M')
      9 # Calculate average monthly price
     10 monthly_avg = (
     11     df_period.groupby('year_month')['valeur_fonciere']
     12     .mean()
     13     .reset_index()
     14     .sort_values(by='year_month')
     15 )

NameError: name 'df_period' is not defined
```

```
[78]: # Average prices by city
price_by_city = df_dvf.groupby('nom_commune')['valeur_fonciere'].mean().
    ↪reset_index()
top_cities = price_by_city.sort_values('valeur_fonciere', ascending=False).
    ↪head(20)

plt.figure(figsize=(10,6))
sns.barplot(data=top_cities, y='nom_commune', x='valeur_fonciere',
    ↪palette='viridis')
plt.title("Top 20 municipalities by average sale price")
plt.xlabel("Average price")
plt.ylabel("City")
plt.show()
```



2.1.8 Evolution of average monthly prices (2019–2024)

The chart shows how the average property sale price in France has evolved month by month concerning the dataset period 2019 and 2024. Despite several sharp spikes likely caused by exceptional transactions or outliers, the overall trend appears relatively stable. This suggests that, while individual high-value sales occasionally distort the monthly average, the broader market did not experience major long-term volatility during this period.

```
[8]: # Convert the date column to datetime format
df_dvf['date_mutation'] = pd.to_datetime(df_dvf['date_mutation'],
    ↪errors='coerce')

# Filter dataset for the years 2019-2024
df_period = df_dvf[df_dvf['date_mutation'].dt.year.between(2019, 2024)]

# Add Year-Month column for monthly aggregation
df_period['year_month'] = df_period['date_mutation'].dt.to_period('M')

# Calculate average sale price per month
monthly_avg = (
    df_period.groupby('year_month')['valeur_fonciere']
    .mean()
    .reset_index()
    .sort_values(by='year_month')
)

# Convert the period to datetime for proper time formatting
```

```

monthly_avg['year_month'] = monthly_avg['year_month'].dt.to_timestamp()


# Line plot showing temporal evolution of average prices
plt.figure(figsize=(12,6))
sns.lineplot(
    data=monthly_avg,
    x='year_month', y='valeur_fonciere',
    marker='o', color='darkorange'
)

plt.title('Evolution of average monthly prices', fontsize=14)
plt.xlabel('Month')
plt.ylabel('Average price (€)')

# X-axis: one major tick per year and minor ticks every 3 months
ax = plt.gca()
ax.xaxis.set_major_locator(mdates.YearLocator())
ax.xaxis.set_minor_locator(mdates.MonthLocator(interval=3))
ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))

plt.xticks(rotation=45)
plt.grid(alpha=0.4)
plt.tight_layout()
plt.show()

```




2.1.9 Number of transactions per month

The chart highlights the seasonality of real estate transactions in France. Sales activity tends to increase during the summer months, peaking in June and July, and again in December, likely due

to year-end transactions and fiscal timing. Conversely, August shows a notable drop, reflecting the slowdown typical of the vacation period.

```
[9]: # Extract the month from the transaction date and count how many sales occur in
      ↪ each month
df_dvf['month'] = df_dvf['date_mutation'].dt.month
transactions_per_month = df_dvf.groupby('month').size().
      ↪ reset_index(name='count')

# Bar plot showing the number of transactions by month
sns.barplot(
    data=transactions_per_month,
    x='month',
    y='count',
    palette='Blues_d'
)
plt.title("Number of transactions per month")
plt.xlabel("Month")
plt.ylabel("Number of sales")
plt.show()
```



2.2 Conclusion

This preliminary exploration of the DVF dataset provides a general understanding of its structure, data quality, and main patterns before conducting a formal Exploratory Data Analysis (EDA).

2.2.1 1. Dataset Structure and Quality

- The dataset is extremely large (~20 million records, 40 columns) and contains both administrative and transactional variables.
- Several cadastral-related columns show over 90% missing values and can be safely removed.
- The core variables relevant for predictive modeling — `valeur_fonciere`, `surface_reelle_bati`, `nombre_pieces_principales`, `type_local`, `date_mutation`, `longitude`, and `latitude` — are complete and consistent.

2.2.2 2. Property Value Distribution

- Property sale prices are heavily right-skewed, with a few extremely high values.
- A logarithmic transformation (\log_{10}) normalizes the distribution and improves interpretability.
- This confirms that future models should work on a log-transformed price variable.

2.2.3 3. Property Type Differences

- “Appartements” represent the majority of transactions, while “Maisons” tend to have higher median prices and larger variability.
- Property type is therefore a key categorical predictor to include in the Machine Learning models.

2.2.4 4. Relationship Between Area and Value

- A clear positive but sublinear correlation exists between built area and property value.
- Some small but high-priced properties likely correspond to urban premium locations.
- The price-per-square-meter ratio will be an important engineered feature.

2.2.5 5. Temporal Patterns

- Transaction volumes exhibit strong annual seasonality, with peaks in spring and early summer.
- Average monthly prices increase from 2019 to 2022, then stabilize slightly after 2023.
- Temporal variables should capture both trend and seasonality components.

2.2.6 6. Geographic Distribution

- Sales are concentrated in Île-de-France, the southeast coast, and major urban centers like the capital (Paris).
- Significant spatial heterogeneity is visible in both transaction density and price levels.
- Geographic variables (latitude, longitude, department) play a central role in property cost.

2.2.7 Summary

The DVF dataset provides a rich and reliable foundation for predictive modeling of property sale prices.

Despite some missing administrative fields, the essential features are clean and informative.

The observed temporal, spatial, and structural heterogeneity indicates that future models should combine log-transformed target values, spatial features, and time-based variables to capture the complex dynamics of the French real estate market.